

Introduction to BitTorrent

Arvid Norberg

arvid@cs.umu.se

<http://libtorrent.net>

Bittorrent


- Introduction
- Efficiency & Reliability
- The incentive mechanism
- Trackerless with DHT

Introduction

- Bittorrent is a system for efficient and scalable replication of large amounts of **static** data
 - Scalable - the throughput increases with the number of downloaders
 - Efficient - it utilises a large amount of available network bandwidth

Introduction

- The file to be distributed is split up in *pieces* and an SHA-1 hash is calculated for each piece



0 1 2 . . .




Introduction

- A *metadata* file (.torrent) is distributed to all peers
 - Usually via HTTP
- The metadata contains:
 - The SHA-1 hashes of all pieces
 - A mapping of the pieces to files
 - A *tracker* reference

Introduction


- The tracker is a central server keeping a list of all peers participating in the *swarm*
- A swarm is the set of peers that are participating in distributing the same files
- A peer joins a swarm by asking the tracker for a peer list and connects to those peers

Introduction



Introduction

Tracker





Goals

- Efficiency
 - Fast downloads
- Reliability
 - Tolerant to dropping peers
 - Ability to verify data integrity (SHA-1 hashes)

Efficiency

- Ability to download from many peers yields fast downloads
- Minimise piece overlap among peers to allow each peer to exchange pieces with as many other peers as possible

Piece overlap



- Small overlap

- Every peer can exchange pieces with all other peers
- The bandwidth can be well utilised

- Big overlap

- Only a few peers can exchange pieces
- The bandwidth is under utilised

Piece overlap


- To minimise piece overlap:
 - Download random pieces
 - Prioritise the rarest pieces, aiming towards uniform piece distribution

Reliability


- Be tolerant against dropping peers
 - Each dropped peer means decreased piece availability
- Maximise piece redundancy
 - Maximise the number of *distributed copies*

Distributed copies

- The number of distributed copies is the number of copies of the **rarest** piece
e.g.



Distributed copies = 2



Distributed copies = 1

Distributed copies



- To maximise the distributed copies, maximise the availability of the rarest pieces
- To increase the availability of a piece, download it
- To maximise the distributed copies:
 - Download the rarest pieces first

Rarest first

- The piece picking algorithm used in Bittorrent is called *rarest first*
- Picks a **random** piece from the set of **rarest** pieces
- No peer has global knowledge of piece availability, it is approximated by the availability among neighbours

Rarest first

- Pick a **random** piece from the set of **rarest** pieces {2, 3}
- Ignore pieces that we already have



The incentive to share

- All peer connections are symmetric
- Both peers have an interest of exchanging data
- Peers may prefer to upload to peers from whom they can download
 - Leads to slow starts
 - Fixed in a recent extension

The incentive to share

- There is a loose connection between upload and download speed
- Each peer has an incentive to upload

Trackerless torrents

- Common problems with trackers
 - Single point of failure
 - Bandwidth bottleneck for publishers
- Solutions
 - Multiple trackers
 - UDP trackers
 - DHT tracker

DHT distributed hash table

- Works as a hash table with sha1-hashes as keys
- The key is the *info-hash*, the hash of the metadata. It uniquely identifies a torrent
- The data is a peer list of the peers in the swarm


DHT distributed hash table

- Each node is assigned an ID
 - in the key space (160 bit numbers)
- Nodes order themselves in a defined topography
 - Makes it possible to search for Ids by traversing the node topography
- Bittorrent uses *kademlia* as DHT

Kademlia bootstrap

- Each node bootstraps by looking for its own ID
 - The search is done recursively until no closer nodes can be found
 - The nodes passed on the way are stored in the routing table
 - The routing table have more room for *close* nodes than distant nodes

Kademlia routing table



- Each node knows much more about close nodes than distant nodes
 - The key space each bucket represents is growing with the power of 2 with the distance
 - Querying a node for a specific ID will on average halve the distance to the target ID each step


Kademlia routing table

- The distance metric is defined as XOR
 - In practice, the distance is 2 to the power of the inverse of the size of the common bit prefix

10011011001	1101010110001
10011011001	0101110101100

Common prefix = 11 Distance $\geq 2^{13}$

Kademlia routing table



Kademlia search

- Each search step increases the common bit prefix by at least one
 - Search complexity: $O(\log n)$

Kademlia distributed tracker

- Each peer *announces* itself with the distributed tracker
 - by looking up the 8 nodes closest to the info-hash of the torrent
 - And send an announce message to them
 - Those 8 nodes will then add the announcing peer to the peer list stored at that info-hash

Kademlia distributed tracker

- A peer joins a torrent by looking up the peer list at a specific info-hash
 - Like a search but nodes return the peer list if they have it

Kademlia distributed tracker

- 8 nodes is considered enough to minimise the probability that all of them will drop from the network within the announce interval
 - Each announce looks up new nodes, in case nodes have joined the network with Ids closer to the info-hash than a previous node

